Paper Title

Abstract

Our abstract.

1 Introduction

Introduction here.

2 Neural Network Background

The human brain is composed of billions of neurons: interconnected electrically excitable cells that form the basis of our intelligence. Each neuron has synapses that receive electrical signals from multiple other neurons. If the sum of these inputs is greater than a certain value, the neuron fires, generating a voltage at its axon. The axon, or the output of the neuron cell, is itself connected to a synapse of another neuron. The interconnections of these neurons form a massive network, where a huge number of inputs are processed in parallel to a set of outputs. The basis of artificial neural networks is to simulate the mathematical properties of these neurons in order to perform similar tasks of mass parallel computation.

2.1 Single Artificial Neuron

Figure 1: Single Neuron Diagram

An artificial neuron simply a mathematical model of a biological neuron, and therefore its functionality is very similar. Each neuron has multiple inputs, each with an individual weight, and one output. The weight of an input is simply a positive or negative fraction that govern the impact of each input on the single output. An artificial neuron's set of weights determines its function, and the function of each neuron in a neural network (in addition to the arrangement of the network) determines the function of the network as a whole. This will

become important as we discuss training Fido's neural network, however for the moment we will adjust our focus to the output of an artificial neural network. By summing each input multiplied by its individual weight we reach a value called the activation, expressed as such mathematically for each input x and weight w:

$$a = \sum_{i=0}^{i=n} x_i w_i \,.$$

(1)

If this activation value is less than a certain threshold, the output is zero. If it is greater than the threshold, the output is one. This activation function most closely resembles the biological model of a neuron, a binary step function. However, a binary output can be somewhat limiting for many applications of neural networks. For example, many of Fido's outputs are gradient rather than linear; as an example, an LED can be given a brightness. For this purpose alternate activation functions can be used with gradient outputs. One such function is a sigmoid function, expressed as such:

$$O(a) = \frac{1}{1 + e^{-\frac{a}{p}}} \,,$$

(2)

for each output O, activation a, and constant p. The sigmoid activation function can also be graphed as below.


Figure 2: Sigmoid Function Graph

This provides us with a gradient output, however output is still limited to positive val-

2

ues. An alternative activation function which allows outputs ranging from -1 to +1 is the hyperbolic tangent activation function:

$$O(a) = \tanh(a)\,. \tag{3}$$

The hyperbolic tangent activation can be graphed as such, demonstrating its greater range and gradient output.


Figure 3: Hyperbolic Tangent Function Graph

2.2 Feedforward Neural Network

A traditional method of arranging artificial neurons in a neural network is called a feedforward network. Neurons are connected as previously described: the output of each neuron is connected to one of the inputs of another neuron. These neurons are organized into layers, as described in the following figure:

Initial inputs are first sent into the input layer. Next, the output of each input layer neuron is sent into each neuron in the hidden layer. The hidden layer processes the inputs to outputs using an activation function and a weight value for each input. There can be any number of hidden layers in a neural network, depending on the complexity of the computation being performed. Finally the outputs from the last hidden layer neurons are sent into each neuron of the output layer, where the final outputs are processed. A concrete example of a

Figure 4: Single Output Feedforward Network

feedforward network is that of Fido. Sensor inputs such as light and sound are sent into the input layer. After the signals have passed through to the output layer, the outputs from the output layer are sent to outputs such as motors, LEDs, and buzzers. The purpose of training a neural network is to make input values correspond to the correct output values, depending on the desired behavior. An example could be making Fido drive when light is applied. The way to give a neural network desired behavior is by controlling the weights of each neuron using a learning algorithm.

3   Learning Implementation

4   Simulation

The robot model chosen for simulation was modeled after easily producible robots on the same scale. The software driving the simulation was intended to be portable enough to work on a hardware implementation, and the model facilitates this goal as well. Additionally, the robot model would have to be easily trainable and debuggable when implemented in hardware; use of a Geiger counter as an input would be unfavorable. Lastly the sensors and design chosen had to facilitate the concept of natural learning, modeling after nature to some degree.

Figure 5: Screenshot of Fido Simulator GUI

4.1   Robot Inputs and Outputs

Multiple inputs were modeled for simulation with outlets for control both by a human operator using sliders and by programmed handlers using a bridge class. A microphone and light sensor were chosen as clear, human modifiable inputs that model after nature and could easily be used for reinforcement training. An infrared light sensor was added as another easily controller variable in a testing setup: a human operator could easily bring closer and farther an IR LED for purposes of training. Additionally sensors for battery level, three axes of accelerometers, and three axes of gyroscopes were added as more complex inputs for Fido to master. The last sensor added was a three axis radio receiver. The purpose of the receiver was to allow location based training of the robot relative to a radio beacon, such as following the beacon or avoiding it. The radio sensor is a stand-in for Bluetooth, Wi-Fi, or any other radio technology; it is common practice in many areas of robotics to use radio beacons for localization.

Fido's outputs were chosen similarly. Two motors allow a drive mechanism known as

5

differential drive, which will be discussed further in the following section. A buzzer of varying
tone and frequency and a multicolor LED complete the set of outputs.

4.2   Implementation and Kinematics

A graphical user interface was made using the SFML multimedia library for C++. Sliders
manually adjust inputs such as accelerometer and gyroscope axes, battery life, and temper-
ature. A colored circle displays the output of the multicolor LED, while the frequency and
volume of the piezoelectric speaker are displayed on the bottom bar. Initially vectors of
motor values were displayed graphically in the top right of the window. This made sense
for initial testing purposes: both competition entrants are experienced with differential drive
robots, and can easily visualize robot movement from robot vectors. However, as we decided
to do more development on the simulator we wanted to allow more complex training on the
simulator, such as path following. Such training requires a visual kinematic simulation of the
robot model.

Figure 6: Differential Drive Kinematics Diagram

Fido's two motors are arranged in a differential drive arrangement. Driving the left motor
acts as a force vector on the left side of the robot, creating a torque around the right wheel.
The same applies with the right wheel. When both motors are activated together, the motor
drives straight. Each motor has a value ranging from -100 to 100, where -100 is full power
backwards, 0 is stopped, and 100 is full power forwards. In order to transform motor values
into a plottable transformation, we must first model the movement that our robot will take.

6

Every movement the robot takes can be interpreted as a rotation of radius R around the ICC, or Instantaneous Center of Curvature. If the robot is moving straight, this radius is simply infinite. The robot travels at an angular velocity ω around the circle, and at any given moment is at the coordinates (x, y) and orientation θ. The length of the robot from wheel-center to wheel-center is defined as l, while the velocity vectors of each motor are defined as V_l and V_r respectively. The passing of time from the last call of motor values is defined as Δt. We can solve for R at any point using the following equation:

$$R = \frac{l}{2} \times \frac{V_l + V_r}{V_r - V_l}. \tag{4}$$

We can solve for ω using the following equation:

$$\omega = \frac{V_r - V_l}{l} \tag{5}$$

And the ICC location using the following equation:

$$ICC = \left[x - R\sin\theta, \quad y + R\cos\theta \right] \tag{6}$$

We can then use these values to solve for the robot's new position, defined as coordinates (x', y') and orientation θ'.

$$\begin{bmatrix} x' \\ y' \\ \theta' \end{bmatrix} = \begin{bmatrix} \cos(\omega\Delta t) & -\sin(\omega\Delta t) & 0 \\ \sin(\omega\Delta t) & \cos(\omega\Delta t) & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x - ICC_x \\ y - ICC_y \\ \theta \end{bmatrix} + \begin{bmatrix} ICC_x \\ ICC_y \\ \omega\Delta t \end{bmatrix} \tag{7}$$

A brief inspection of these equations verify their performance in certain use cases. If $V_r = -V_l$ R becomes zero, as the robot turns around it's center. If $V_r = V_l$ R becomes infinite, as the robot is traveling in a straight line. However the methodology by which the equation simplifies in the case of $V_r = V_l$ involves division of zero, which can be problematic

in computer programming. Therefore we must first check if $V_r = V_l$, and if so substitute an alternate equation as a simplification:

$$\begin{bmatrix} x' \\ y' \\ \theta' \end{bmatrix} = \begin{bmatrix} x + V_l \Delta t \cos \Theta \\ y + V_l \Delta t \sin \Theta \\ \theta \end{bmatrix} \tag{8}$$

The use of V_l in this equation rather than V_r is unimportant, as the values are equal. Using these equations were were able to fully simulate Fido's kinematic model, and attach simulator motor inputs to Fido's outputs to visualize learning taking place. The black rectangle in the upper right corner of the simulator is the robot, having been moved as part of training. The red dot near the rectangle is a graphical representation of a radio beacon. As adjusting sliders to represent the location of a radio beacon relative to the robot would be impractical, we decided to implement a beacon that could be placed and dragged by right clicking on the simulator. Simulated sensor readings of beacon strength on two axis are gathered using an inverse square law, as applies to radio waves in general. These readings are then displayed in the sliders and can be manually altered as well. The radio beacon can be removed by a human operator by pressing the P key. This has been especially helpful in the task of training Fido to follow a radio beacon.

5   Results

Results, testing, and applications go here.

5.1   Training Methods

5.2   Findings

5.3   Further Applications

6   Discussion

7   Conclusion

Restate, discuss further study, improving experimentation, etc.

References

Dini, S., & Serrano, M. (2012). Combining q-learning with artificial neural networks in an adaptive light seeking robot.

Dudek, G., & Jenkin, M. (2000). *Computational principles of mobile robotics.* New York, NY, USA: Cambridge University Press.

Gaskett, C., Wettergreen, D., & Zelinsky, A. (1999). Q-learning in continuous state and action spaces. In *Australian joint conference on artificial intelligence* (pp. 417–428).

Kim, D. S., & Papagelis, A. J. (n.d.). *Multi-layer perceptron: Artificial neural networks.* `http://www.cse.unsw.edu.au/ cs9417ml/MLP2/`.

MacLeod, C. (2010). *An introduction to practical neural networks and genetic algorithms for engineers and scientists* (Tech. Rep.). Robert Gordon University.

Werbos, P. (1974). Beyond regression: New tools for prediction and analysis in the behavioral sciences.